Where Guard Oversight Fails: Input Length Degrades Error Detection and Panel Agreement Together

Draft. Every number traced to `results/tables_repriced/`*; see* `docs/findings.md` *for the table-by-table provenance.*

Abstract

Content-moderation guard models make mistakes, and deployments cover for them by routing uncertain decisions to human reviewers. Which decisions to route, and whether the compute spent identifying them would be better spent on a larger guard, is an allocation question that has not been priced. We evaluate three generative guards (Llama-Guard-3-1B/8B, Qwen3Guard-Gen-8B) on five moderation datasets, 26,958 predictions in a single hardware environment, under a pre-registered protocol fixed before any data was collected. Policies are compared on residual risk – errors reaching users after review – against two cost axes that we deliberately never convert into one another.

Three results. First, the best policy depends on the review budget: a single 8B guard on its own confidence wins at 1-5% and 50% review, while a three-guard agreement panel wins only at 10-20%, at 2.5x the compute. Second, there is a large and precisely estimated *identification gap*: routing each prompt to the guard that will answer it correctly would reduce error 26.3% (95% CI [23.6, 29.2]) at half the compute of the 8B alone, and the best available escalation signal captures none of it (realised gain 95% CI [-0.0017, +0.0012]). Third, and we argue most consequentially, both failures share an axis. The 2.85% of prompts on which all three guards agree *and are wrong* are confidently wrong, are absent entirely from short curated benchmarks, and grow roughly ninefold from the shortest to the longest length quintile; the same length variable destroys the behavioural perturbation signal (rank correlation -1.0 between mean prompt length and usable variant count).

On long inputs the signals that would route a case for review stop working exactly where the panel stops catching it. Short curated benchmarks cannot observe this, which we suggest is why it has not been reported.

1. Introduction

A guard model sits in front of a deployed language model and decides whether an incoming message should be blocked. It is wrong in both directions: it blocks benign requests and admits harmful ones. Operators compensate with human review, but review is scarce – a realistic moderation team inspects on the order of 1-20% of decisions, not half.

That scarcity makes the operator's problem an allocation problem. A fixed oversight budget can buy a larger and more accurate guard, additional compute spent checking the guard's answer, or more reviewer hours. These compete, and the exchange rate between them is unknown.

The nearest prior work, Safety-Flag (arXiv:2609.19072), establishes that a guard's own token-logprob confidence ranks its errors usefully, across seven benchmarks and thirteen models. It contains no cost analysis of any kind, and reports risk only at 50%, 80% and 100% coverage on a protocol balanced ~50/50 by construction. The regime an operator inhabits – 1-20% review at a 1-7% base rate – is untouched.

This paper occupies that regime and prices it. Our contributions:

1. **An oversight frontier.** Every (guard, routing signal, review budget) policy priced on two axes that are never converted into one another, over 26,958 predictions in one hardware environment (§3).
2. **Budget-dependent policy selection** (§5.2). The winner changes with the review budget, non-monotonically.
3. **The identification gap** (§5.4). We separate what routing *could* buy from what available signals *do* buy, with bootstrap intervals. The gap is large and the capture is statistically zero.
4. **A length axis for guard-oversight failure** (§5.5). Unanimous panel failure and behavioural-signal collapse are both governed by input length, which short curated benchmarks hold fixed.

Pre-registration. The hypotheses, decision thresholds and gate structure were committed to a public document while the results directory was empty, and frozen when the first real predictions existed. Two consequences are visible in this paper and we report both: our principal hypothesis – that a cheap guard with expensive checking beats an expensive guard with free checking – was *not* supported, and a mechanism we proposed post-hoc was specified, tested on eight

panels, and rejected (§5.3, §7).

---

2. Related work

Selective prediction for moderation. Safety-Flag (arXiv:2609.19072) benchmarks reliability and calibration of content moderators across seven datasets and thirteen models, reporting AURC as its primary instrument. It establishes that token-logprob confidence is usable and that temperature scaling improves calibration numbers without changing prediction ordering. We replicate the first finding (§5.1) and our §5.1 result on the logit margin is the second in a sharper form.

Its Appendix D.1 is a direct headwind for behavioural signals: sampled-answer agreement over five stochastic generations loses to token-logprob confidence on all six models tested. We differ in varying the *input* rather than the *decoder* – a guard scored greedily on one label token has no sampling variance at all – but our §5.5 result is, in outcome, consistent with theirs.

Perturbation and instability. arXiv:2402.13006 finds that examples which flip under perturbation already carry higher baseline uncertainty, raising the possibility that instability is a costly restatement of a small margin. Our gate S4 was constructed to detect precisely this; we could not evaluate it on realistic inputs, for reasons that became a result (§5.5). Semantic entropy and semantic-invariant perturbation sampling (Kuhn & Gal and successors) establish the technique for hallucination detection; its application to guard error routing is what is new, not the method. arXiv:2603.21172 argues entropy alone is insufficient for safe selective prediction, supporting a search beyond the obvious signal.

Polarized guard scores. An industry benchmark reports verdict probabilities "almost perfectly polarized, with 99.8% of scores pinned at zero or one" for some guards. We measured this on our own models rather than assuming it, and it does not hold: 12-37% pinned, with usable interior resolution throughout (§5.1). Safety-Flag likewise found dedicated guards better calibrated than general-purpose models.

Unoccupied. Across this literature we find no cost analysis of guard oversight, no treatment of the 1-20% review regime, no cross-model disagreement as a routing signal, and no comparison of logit margin against probability.

---

3. Method

3.1 Policies and the objective

A **policy** is a (guard, routing signal) pair together with a review budget `b`. The signal assigns every prompt a score; the `b` least-confident fraction goes to a human, who is assumed to correct any error shown to them. No thresholds are involved – only the ordering of a signal is ever used, so signals never need to be calibrated onto a common scale.

The objective is **residual risk**:

```
residual = errors the guard made  -  errors the reviewer saw
```

We use residual rather than accuracy or errors-caught because both flatter a weak guard. A guard making 1,614 errors and catching 1,128 ships more harm than one making 943 and catching 561.

3.2 Two cost units, never mixed

Compute and reviewer hours are not denominated in the same thing, and any exchange rate dates badly. Nothing in this study converts one into the other. Policies are reported on a two-dimensional frontier and the reader supplies their own prices.

Compute is measured latency per decision on the evaluation machine, normalised. A panel's price is the **sum of its members' costs** – it runs every member on every prompt, so its price is a property of the panel, not of whichever member is nominated to decide.

3.3 Guards, data, signals

Three generative guards, scored at fp16 through llama.cpp:

guard	error rate	AUROC	flag rate	compute
Llama-Guard-3-8B	0.1049	0.9226	0.200	0.319

guard	error rate	AUROC	flag rate	compute
Qwen3Guard-Gen-8B	0.1157	0.9615	0.314	0.372
Llama-Guard-3-1B	0.1796	0.8704	0.283	0.102
(three-guard panel)				*0.793*

Five datasets, 8,986 prompts after normalisation, pooled base rate 22.5%: ToxicChat (4,972; 7.1% harmful; real user traffic), WildGuardTest (1,699), OpenAI-Moderation (1,665), XSTest (450; benign-but-borderline), HarmBench (200; all harmful).

Signals: native confidence `|p_unsafe - 0.5|` (free, the baseline), logit margin `|margin|` (free), three-guard modal agreement (3x compute), and perturbation instability over six meaning-preserving rewordings (6x compute).

Environment. All 26,958 predictions were scored on one NVIDIA L40S (`env_hash 6ccc83087810`). Every prediction row carries this hash and the analysis refuses to treat latency as comparable across differing hashes. Compute prices are therefore valid within this environment and only within it.

---

4. Pre-registered gates

Each gate can fail, and was verified against synthetic data whose answer is known by construction, in both directions, before any real scoring (13/13 checks). Thresholds were fixed in advance: a 5% relative residual reduction for the allocation gate, 0.01 AURC for signal comparisons.

Gate	Question	Result
A1	At matched compute, does anything beat the default?	`CHEAPER_GUARD_WINS` (see §5.6)
A2	Does the winner change with budget?	`BUDGET_DEPENDENT`
S1	Is native confidence usable?	`NATIVE_USABLE` (3/3)
S2	Does the margin beat the probability?	`NO_DIFFERENCE`
S3	Does any alternative beat native confidence?	`ALTERNATIVE_BEATS_NATIVE` (1/3 guards)
S4	Is instability independent of the margin?	`NOT_EVALUABLE` (§5.5)
S5	Does guard *kind* decide usability?	`NOT_EVALUABLE` – no aligned encoder

---

5. Results

5.1 Native confidence is usable, and the margin is exactly the probability

All three guards have usable native confidence (deferral efficiency 0.667-0.745, tie fractions below 0.003), replicating Safety-Flag on models it did not test.

The logit margin – free to compute and hypothesised to carry ranking information destroyed by sigmoid saturation – is *identical* to the probability. AUROC and AUROC-on-margin agree to seven decimal places on all three guards (delta AURC <= 1.8e-08). The saturation never occurred: median `|margin|` is 4.0-5.8, far short of the ~37 at which float64 sigmoid pins to 1.0, and only 12-37% of scores are pinned at all.

The premise motivating behavioural signals – that guard scores are degenerate – does not hold for these guards. It should be measured rather than assumed.

5.2 The best policy depends on the review budget

Review budget	Winning policy	Residual	Compute
1%	Llama-Guard-3-8B + native	0.0999	0.319
2%	Llama-Guard-3-8B + native	0.0963	0.319
5%	Llama-Guard-3-8B + native	0.0832	0.319
10%	**Qwen3Guard-8B + panel agreement**	**0.0634**	**0.793**

Review budget	Winning policy	Residual	Compute
20%	**Qwen3Guard-8B + panel agreement**	**0.0376**	**0.793**
50%	Llama-Guard-3-8B + native	0.0112	0.319

The expensive signal wins in a **middle band only**, at 2.5x the compute. The pattern is not monotone in budget, and §7 reports a mechanism for the band that we proposed and then rejected.

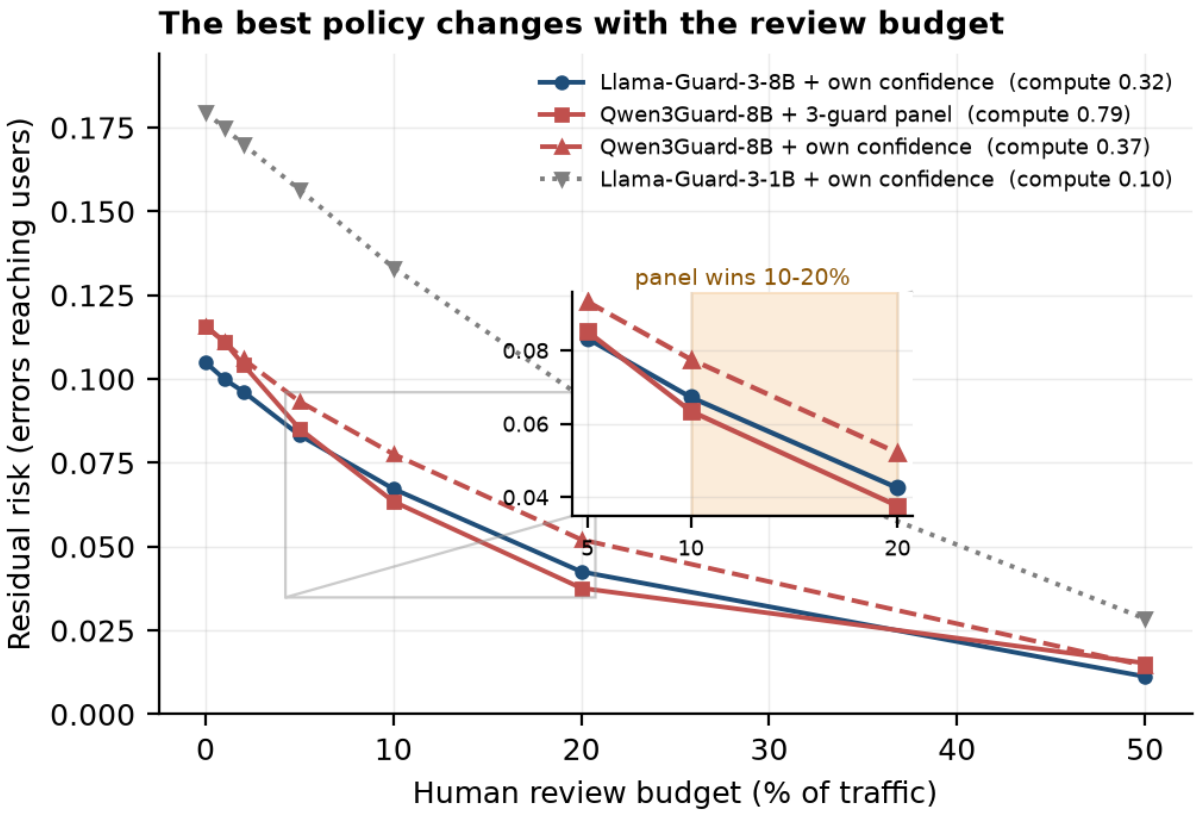


Figure 1. Residual risk against human review budget for four policies. The 8B guard on its own confidence wins below 10% review and at 50%; the three-guard panel wins only in the 10-20% band, at 2.5x the compute. Inset magnifies the crossover. The 1B guard never wins at any budget.

The 1B guard never wins at any budget. Our principal pre-registered hypothesis is not supported.

5.3 Panel agreement is two-valued and helps only the weakest guard

With three guards voting, modal agreement takes two values. The tie fraction is 1.0000 on every guard and every dataset: the signal partitions traffic into a flagged block and an unflagged one and imposes no order within either.

The flagged block is 23.8% of traffic (2,141 prompts) and contains 72.9-84.1% of each guard's errors. But the benefit is not uniform:

guard	delta AURC vs native	95% CI	Holm p
Llama-Guard-3-1B	**-0.0142**	[-0.0201, -0.0085]	0.000
Llama-Guard-3-8B	+0.0158	[+0.0101, +0.0204]	1.000
Qwen3Guard-8B	+0.0092	[-0.0011, +0.0103]	1.000

(Negative is better.) Agreement is a large, correction-surviving improvement for the 1B and runs in the *wrong* direction for both 8B guards. A strong guard gains nothing from consulting a weaker one. We note this because "ensemble disagreement helps" is commonly asserted without this split.

5.4 The identification gap

Every policy above has one guard decide every prompt. We additionally evaluate a cascade – the 1B decides what it can, escalating low-confidence cases to the 8B – recomputed from the existing predictions with no new inference.

policy	error	95% CI	compute
8B alone	0.1049	[0.0987, 0.1110]	0.319
cascade, escalating on 1B confidence	0.1052	[0.0989, 0.1112]	0.294
cascade, oracle escalation	**0.0773**	**[0.0720, 0.0827]**	**0.159**

Bootstrap over prompts, 2,000 resamples:

- **Realised** gain (8B - cascade): **-0.0002, CI [-0.0017, +0.0012]** – straddles zero.
- **Available** gain (8B - oracle): **+0.0276, CI [+0.0243, +0.0309]** – excludes zero.

The two guards fail on sufficiently different cases that perfect routing between them would cut error **26.3% [23.6%, 29.2%]** at **half** the 8B's compute. The best available escalation signal captures none of this.

The bottleneck is not compute, which is cheap and was spent freely here. It is the inability to identify which decisions are wrong.

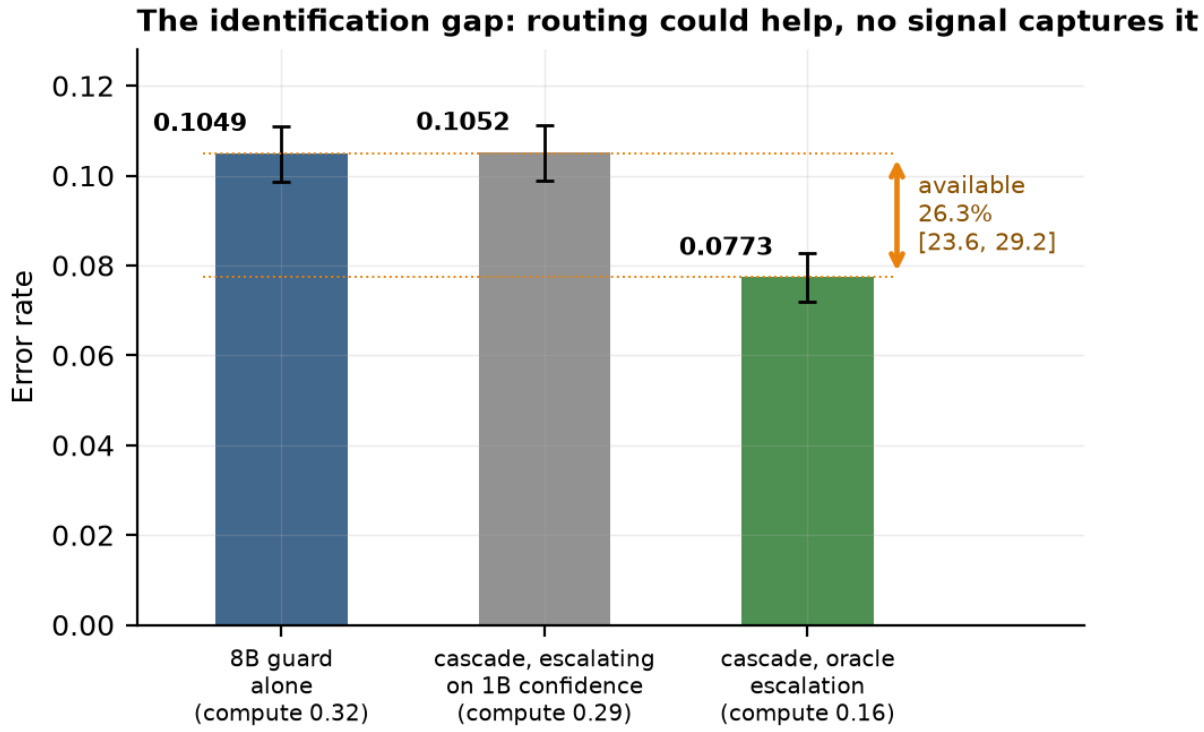


Figure 2. The identification gap. Error bars are 95% bootstrap intervals over 2,000 resamples. Escalating on the 1B guard's own confidence is indistinguishable from the 8B alone. Escalating on an oracle cuts error 26.3% at half the compute. Everything between the dotted lines is available and uncaptured.

5.5 Where the panel is blind, and why length is the axis

On 2.85% of prompts (256) all three guards agree **and all three are wrong**: 167 benign prompts flagged, and 89 harmful prompts admitted. Three properties make this set more than a residual.

They are confidently wrong. Median |p - 0.5| is 0.459-0.480 when the panel is unanimously wrong against 0.493-0.499 when it is unanimously right – both near the 0.5 ceiling. No confidence-based routing can reach them, because confidence does not separate them. This is why the 2.85% is a floor rather than a budget shortfall.

Short curated benchmarks cannot see them. HarmBench contributes **zero** unanimous failures. Every one of the 256 comes from realistic or long-form traffic:

dataset	n	unanimously wrong	rate
WildGuardTest	1,699	87	5.12%

dataset	n	unanimously wrong	rate
OpenAI-Moderation	1,665	82	4.92%
XSTest	450	11	2.44%
ToxicChat	4,972	76	1.53%
HarmBench	200	**0**	**0.00%**

Length drives it, within datasets as well as across them. Median prompt length of unanimous failures against the rest of the same dataset: 967 vs 363 characters (OpenAI-Moderation), 548 vs 236 (WildGuardTest), 83 vs 59 (ToxicChat). Pooled by length quintile, the unanimous-failure rate runs 0.65% (median 30 chars), 2.17%, 3.07%, 2.28%, **6.12%** (median 869 chars) – roughly ninefold from shortest to longest.

The same variable governs the behavioural signal. Our perturbation family – meaning-preserving surface edits: inserted greetings, expanded contractions, opener substitutions – produced too few distinguishable variants to score on three of five datasets, and the diversity gate refused them before inference:

dataset	mean prompt chars	usable variants (of 6)
XSTest	44	5.0
HarmBench	89	4.4
ToxicChat	167	3.5 (blocked)
WildGuardTest	480	2.4 (blocked)
OpenAI-Moderation	645	2.0 (blocked)

Rank correlation between mean prompt length and usable variant count is **-1.0**. The mechanism is arithmetic: a surface edit is a fixed-size change, so its relative magnitude is that size divided by prompt length. Where it could be measured (XSTest), instability lost outright – AURC 0.0295 against native's 0.0157, catching 10 errors at 5% review where native caught 12, at six times the compute.

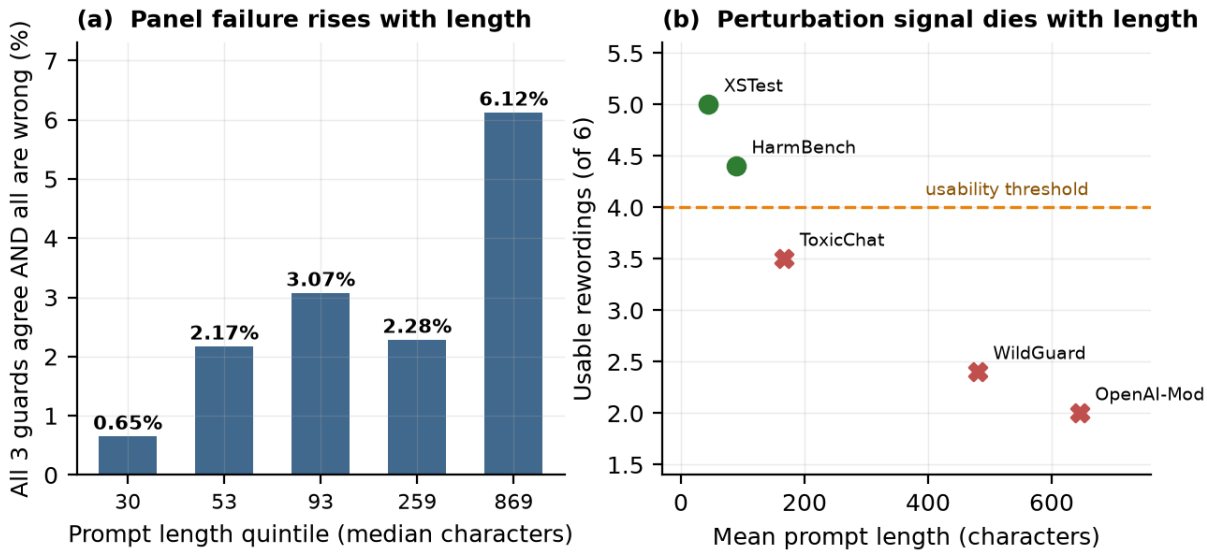


Figure 3. Input length governs both failures. (a) The rate at which all three guards agree and all three are wrong, by prompt length quintile: 0.65% at 30 characters rising to 6.12% at 869. (b) Usable rewordings per prompt against mean prompt length, by dataset; the three datasets below the usability threshold are the three long-form ones, and could not be scored.

> ***Guard oversight degrades with input length from both directions at once.*** *The signals that would route a case for review stop working exactly where the panel stops catching it. This is a property of the inputs, not of a method, and short curated benchmarks hold the variable fixed.*

5.6 Discrimination is not deployment

Gate A1 reports that one guard beats another at matched compute by 10.7%. The comparison is Llama-Guard-3-8B against Qwen3Guard-Gen-8B, *both using native confidence* – there is no difference in selection, and we do not present

this as an allocation result. What it does show is worth separating:

Qwen3Guard has the highest AUROC of the three (0.9615 vs 0.9226) and the worse deployed error rate (0.1157 vs 0.1049), because its operating point over-flags: 31.4% flagged against a 22.5% base rate. **Selecting a guard by AUROC picks Qwen; selecting by deployed residual risk picks Llama-Guard-3-8B.** Guard evaluations predominantly report the former.

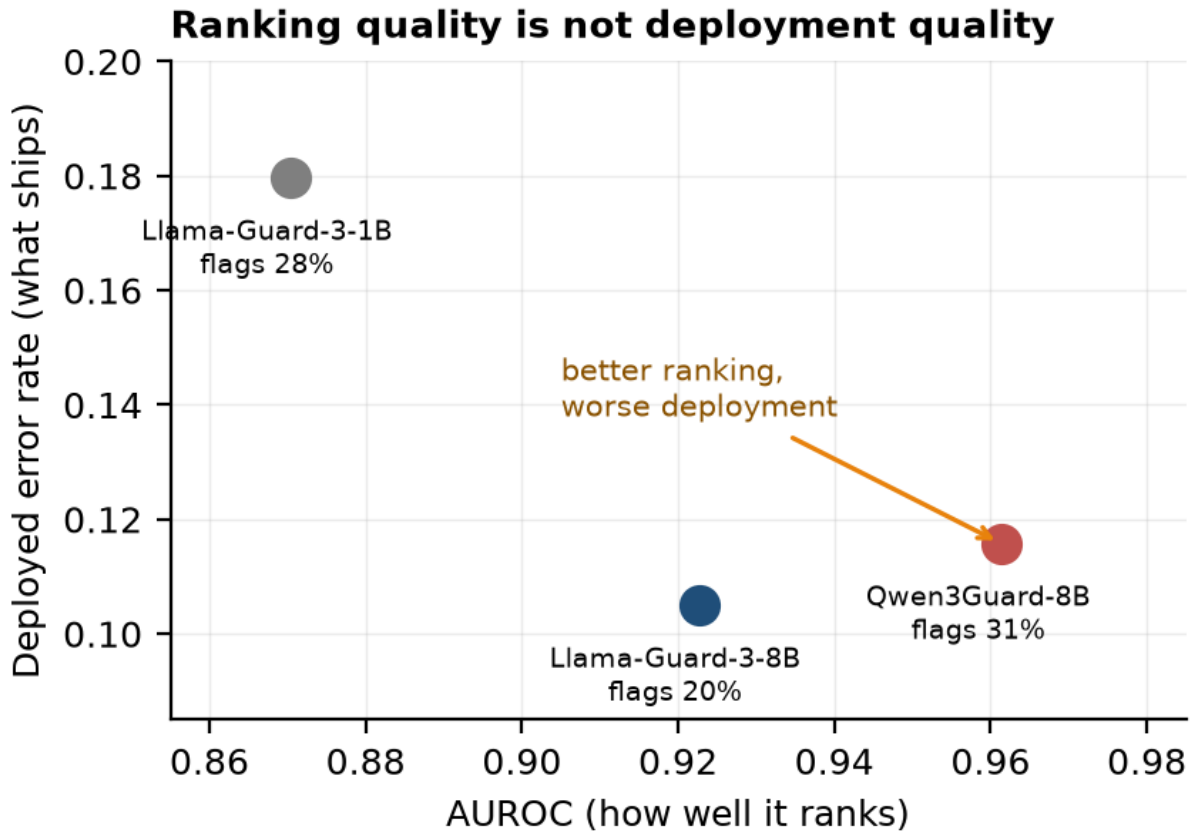


Figure 4. Ranking quality against deployed error. Qwen3Guard-Gen-8B ranks best and deploys worst, because it flags 31.4% of traffic against a 22.5% base rate.

6. Discussion

Five routes to reducing residual risk were priced, and each is bounded:

route	bound
logit margin	identical ordering to the probability
perturbation instability	unmeasurable at realistic prompt length
cross-model disagreement	two-valued; wins only in a 10-20% band
cascaded escalation	captures 0 of 0.0276 available headroom
any panel routing	2.85% unanimous-failure floor, confidently wrong

None is a compute limit. Each is a failure to identify which decisions are wrong. The practical recommendation that follows is unglamorous: at review budgets below 10%, run the best single guard on its own confidence and spend nothing on routing, because nothing available beats it.

The length result suggests where effort should go instead. If unanimous failure and signal collapse are both governed by input length, then evaluation on short curated benchmarks systematically understates both, and progress requires length-proportional perturbations – which means a paraphrase model, which reintroduces a confound we deliberately excluded: a paraphraser can sanitise a harmful prompt, making a flipped verdict uninterpretable.

7. A mechanism we proposed and rejected

The natural explanation for the 10-20% band in §5.2 is that the signal is exhausted once the review budget exceeds its flag rate, predicting that the band *ends* at the flag rate. We tested it: three guards give three pairs plus the triple, each panel with its own flag rate, residuals recomputed over a 1-50% grid with ties broken at random over 20 seeds.

The prediction fails. Spearman correlation between flag rate and band end is 0.435 over eight panels, and the band overshoots the flag rate in seven of eight, by 0.159 on average. The flag rate is where agreement stops *gaining*, not where it stops *leading*. We report this rather than omitting it.

---

8. Limitations

1. **Single hardware environment.** Compute prices are valid on one L40S and nowhere else, by design. A1's notion of "the default" is the highest-latency free-signal policy, which is hardware-dependent and could invert elsewhere.
2. **Three guards, one family of alternative signals.** Panel agreement is effectively the only alternative signal we could evaluate; S2 was null by construction, S4 unmeasurable, S5 untested.
3. **No encoder guard.** The one candidate scored AUROC 0.599 against our label – a toxicity head rather than a harmful-request head – and was withdrawn before scoring. S5 remains open and we do not claim it.
4. **Sections 5.4, 5.5 and 7 are exploratory**, derived after seeing the pre-registered results. The oracle in §5.4 bounds this guard pair only.
5. **The length result is correlational.** We show the association holds within datasets, which rules out composition, but not that length is causal rather than a proxy for topic or complexity.
6. **Perturbation was tested in its cheap form only.** We did not test a paraphrase model, for the confound reason above, so §5.5 bounds surface perturbation and not behavioural signals in general.

---

9. Conclusion

Given a fixed oversight budget for a safety guard, at the review rates deployments actually run, the answer is to buy the better guard and spend nothing on routing. Everything we priced that costs more does not beat it outside a narrow band.

The reason is identifiable and quantified: perfect routing between two guards we already have would cut error by a quarter at half the compute, and no available signal captures any of it. And the cases where oversight fails hardest are the long ones – where the panel is confidently and unanimously wrong, and where the signals that would flag them have stopped working. Short curated benchmarks hold that variable fixed and cannot observe either effect.

Reproducibility

Pre-registration, protocol, gate definitions and the full findings-to-table provenance are in `docs/`. Gates are verified against synthetic data with known answers (13/13) and that check is re-run after any change to the analysis code. Claims are licensed by a generated table rather than curated by hand; claims the table does not license are not made here.

References

Annotated for reading. Each entry says **what they did**, **what we found**, and **the difference**. Read-status is from `docs/related_work.md`: only Safety-Flag has been read in full. Do not cite a `SNIPPET` entry as if read — read it first or drop it.

> ***Verify every ID before submission.*** *The 2026-dated arXiv identifiers are transcribed from* `related_work.md`*; the older ones are reconstructed and must be checked against the actual listing.*

A. The direct neighbour — read this first

[1] Safety-Flag: A Unified Benchmark for the Reliability and Calibration of LLM Content Moderators. arXiv:2609.19072. `READ IN FULL 2026-09-28` <https://arxiv.org/abs/2609.19072>

- **What they did.** Seven safety benchmarks (BeaverTails, XSTest, Ethics,

WildGuard, Aegis, ToxiChat, ToxiGen) under one balanced flag/do-not-flag protocol. Thirteen models — six general-purpose LLMs, four dedicated guards (incl. Llama Guard 3 8B), three reference models, **all generative**. Three dimensions: error direction, probability calibration, confidence-based error ranking. AURC is their primary metric. ~1,400 items, ~50/50 harmful/benign. Item-level decisions and confidence scores released.

- **Their finding.** Token-logprob confidence lowers error for every model, and

gives the lowest AURC of the signals they test. Temperature scaling improves calibration numbers **without changing prediction ordering**.

- **Where we agree.** Our S1 replicates them on models they did not test

(deferral efficiency 0.667–0.745, 3/3 usable). Our S2 is their temperature observation in a sharper form: the logit margin is a strictly monotone transform of the probability and ranks identically to seven decimals.

- **The difference, and it is the reason this paper exists.** They have **no

cost analysis of any kind** — no latency, no compute budget, no comparison of spending compute against spending reviewer time. And they report risk only at 50%, 80% and 100% coverage on a 50/50-balanced protocol. Risk@0.5 means a human reviews half of all traffic. We occupy 1–20% review at a 7.1% base rate (ToxicChat), which their design cannot address.

- **Read §4 (confidence signals), §5.5 Table 15 (risk at coverage), and

Appendix D.1 — see [2].**

[2] Safety-Flag, Appendix D.1 — the headwind. Same paper.

- **What they did.** Compared "sampled-answer agreement" (five stochastic

generations per item, scored on verdict agreement) against token-logprob confidence. Token-logprob wins on **all six** models (their Table 14): e.g. Qwen2.5-32B 0.048 vs 0.115, Gemma-2-9B 0.082 vs 0.163, Llama-3.1-8B 0.300 vs 0.427.

- **Why it matters to us.** A reviewer will ask why paraphrase-consistency

should succeed where sample-consistency failed. Our §5.5 outcome is, in direction, **consistent with theirs** — behavioural signals lost — and we say so rather than claiming a reversal.

- **The difference.** They vary the *decoder*; we vary the *input*. A guard

scored greedily on one label token has no sampling variance at all, so their signal is unavailable to us. More importantly, we identify **why** ours failed — a length mechanism (§5.5) they could not observe, since their items are short and balanced.

B. Uncertainty and selective prediction

[3] Investigating the Impact of Model Instability on Explanations and Uncertainty. arXiv:2402.13006. `SNIPPET` <https://arxiv.org/pdf/2402.13006>

- **What they did.** Examples that flip under perturbation already carry higher

baseline uncertainty.

- **The threat.** If instability merely restates a small margin, it costs ~6–10x

inference for nothing. Our gate S4 was built specifically to detect this.

- **The difference.** We never got to test it on realistic inputs. Our

perturbation family was refused by a diversity gate on three of five datasets — and *that refusal became the result* (§5.5). Their hypothesis remains open for length-proportional perturbations.

[4] Entropy Alone is Insufficient for Safe Selective Prediction in LLMs. arXiv:2603.21172. `SNIPPET` <https://arxiv.org/pdf/2603.21172>

- **What they did.** Entropy is a weak uncertainty signal for safe selective

prediction.

- **Relation.** Supports searching beyond the obvious signal, which motivated

our five-signal design. Our conclusion is bleaker: every alternative we could price is bounded, and the binding constraint is error identification (§6).

[5] Uncertainty-Aware Abstention in LLMs with Provable Alignment Guarantees. arXiv:2607.04430. `SNIPPET` <https://arxiv.org/pdf/2607.04430>

- **What they did.** Abstention with provable guarantees.
- **The difference.** For *generation*, not for a moderation classifier, and

with no human-review budget and no cost axis.

[6] Semantic Uncertainty / semantic entropy. Kuhn, Gal & Farquhar, ICLR 2023, arXiv:2302.09664, and successors. `SNIPPET — VERIFY ID`

- **What they did.** Semantic-invariant perturbation sampling for hallucination

detection. **The technique behind our signal 3.**

- **The difference.** The method is established; applying it to guard

error-routing is what is new. We do not claim the technique. We report the condition under which it fails to transfer: input length (§5.5).

C. Guardrail evaluation and operating points

[7] Artificial Analysis guardrail benchmark. Industry, not peer-reviewed. `SNIPPET` <https://artificialanalysis.ai/articles/guardrail-safety-benchmark>

- **What they report.** Some reasoning guards have verdict probabilities

"almost perfectly polarized, with 99.8% of scores pinned at zero or one, leaving no range to tune a threshold," retaining ~10% recall at a 1% FPR budget.

- **We measured it and it does not hold for our guards.** 12–37% pinned, interior

resolution 0.87–0.97, all three `USABLE`. Median |margin| 4.0–5.8 against the ~37 needed for float64 sigmoid saturation.

- **The difference.** Treat this as **motivation, not evidence**. It is the

premise that justifies behavioural signals, and on our models the premise is false. Safety-Flag independently found dedicated guards better calibrated than general-purpose models.

[8] Reasoning's Razor. arXiv:2510.21049. `SNIPPET` <https://arxiv.org/pdf/2510.21049>

- **What they did.** Reasoning improves accuracy but can hurt recall at critical

operating points in safety and hallucination detection.

- **Relation.** Our closest methodological relative — operating-point-aware

evaluation of safety detectors. Their independent variable is reasoning; ours is the oversight budget.

[9] Safety Under Scaffolding. arXiv:2603.10044. `SNIPPET` <https://arxiv.org/pdf/2603.10044>

- **What they did.** Apparent large safety differences turn out to be

measurement artifacts; scores are sensitive to how answers are extracted.

- **Our closest ally in framing.** Same shape of argument, different cause: our

§5.6 shows a guard with strictly better AUROC (0.9615 vs 0.9226) is the worse deployment (0.1157 vs 0.1049) because its operating point over-flags.

[10] CPU-deployable safety classification. arXiv:2608.21570. `SNIPPET` <https://arxiv.org/pdf/2608.21570>

- **What they argue.** A *fixed* decision threshold for cross-model comparison,

explicitly forbidding tuning each model to its own operating point.

- **This is the objection we must answer**, and §5.6 is where. We do not say

fixed thresholds are wrong. We say a fixed-threshold number is a *joint* measurement of discrimination and operating point, and reporting it alone makes the two indistinguishable. We report both. Write this as a paragraph, not a footnote.

[11] When Benchmarks Lie. arXiv:2602.14161. `SNIPPET` <https://arxiv.org/pdf/2602.14161>

- **What they did.** Malicious-prompt classifiers under true distribution shift.
- **Relation.** Adjacent to our §5.5: HarmBench produces **zero** unanimous

failures while long-form sets produce 5%. A benchmark can be silent about a failure mode it holds fixed.

D. The guards

[12] Llama Guard. Inan et al., arXiv:2312.06674; model card for Llama-Guard-3-8B. `SNIPPET — VERIFY ID` <https://huggingface.co/meta-llama/Llama-Guard-3-8B>

- Our reference guard and the size axis (1B). The card reports int8 at F1

0.936–0.939 with "negligible impact" — **at a fixed threshold**, the statistic that cannot separate preserved discrimination from a moved operating point.

[13] Qwen3Guard Technical Report. arXiv:2510.14276. `SNIPPET` <https://arxiv.org/html/2510.14276v1>

- Our replication family. Relevant to §5.6: Qwen3Guard-Gen-8B has the best AUROC

in our panel and the worse deployed error rate.

[14] Llama Guard 3-1B-INT4. arXiv:2411.17713. `SNIPPET` <https://arxiv.org/pdf/2411.17713>

- Pruning plus 4-bit for mobile. Reports improved safety *and* efficiency.

Context for why a 1B guard is a live deployment option — which makes our result that it never wins at any budget worth stating.

E. Datasets — all IDs need verifying

*[15] ToxicChat.** Lin et al., EMNLP Findings 2023, arXiv:2310.17389. *Our headline set: 4,972 rows, 7.1% harmful, real user traffic.*

*[16] HarmBench.** Mazeika et al., ICML 2024, arXiv:2402.04249. *200 rows, all harmful. Produces zero unanimous failures (§5.5) — the key negative control.*

*[17] XSTest.** Röttger et al., NAACL 2024, arXiv:2308.01263. *450 rows, benign-but-borderline. The only set where perturbation was measurable.*

*[18] WildGuard.** Han et al., NeurIPS D&B 2024, arXiv:2406.18495. *1,699 rows. Highest unanimous-failure rate (5.12%) and 70 of our 89 safety-critical false negatives.*

*[19] OpenAI Moderation.** Markov et al., AAAI 2023, arXiv:2208.03274. *1,665 rows. Longest prompts (645 mean chars) and the largest length gap between unanimous failures and the rest (967 vs 363).*

F. Method

*[20] Selective classification / risk–coverage.** El-Yaniv & Wiener, JMLR 2010; Geifman & El-Yaniv, NeurIPS 2017. `VERIFY` *Source of the risk–coverage curve and AURC, our primary instrument — shared with Safety-Flag, which is what makes the comparison checkable.*

*[21] Holm, S. (1979).** A simple sequentially rejective multiple test procedure. Scandinavian Journal of Statistics 6(2), 65–70. *Our multiple-comparison correction in S3. Only the 1B guard's agreement result survives it.*

Reading order

If you read four things, read these:

1. **[1] Safety-Flag in full** — it defines the space you are in.
2. **[2] its Appendix D.1** — the specific result a reviewer will quote at you.
3. **[10] CPU-deployable safety classification** — the methodological objection to how you compare guards.

4. **[9] Safety Under Scaffolding** — the framing to borrow.

[3], [6] and [7] are the three you must read before claiming anything about perturbation or polarization, because each one either predicts your result or supplies the premise you overturned.